

Modern Operating Systems

Chapter 3 – Memory Management

Zhang Yang
Autumn 2022



Contents of the Lecture

■ 3.1 No Memory Abstraction

- Mono program without memory abstraction
 - Overlays

- Multiple programs without memory abstraction

■ 3.2 A Memory Abstraction: Address Spaces

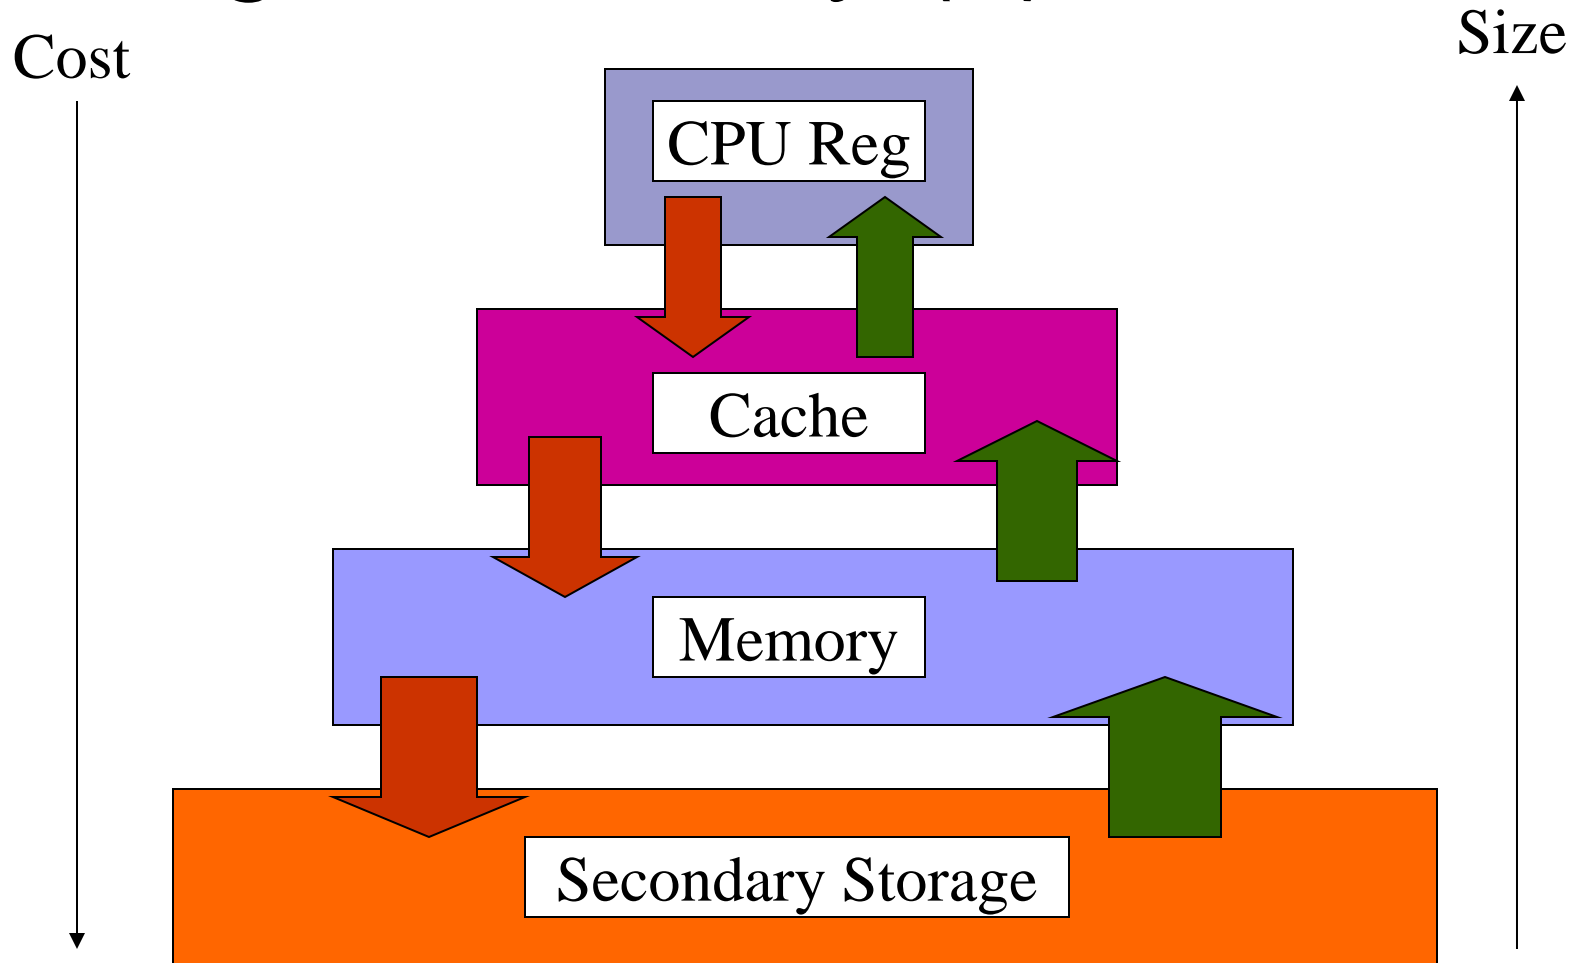
- The notion of an address space
- Dynamic Relocation
- Swapping
- Managing Free Memory
- Storage Placement Strategy



Storage Hierarchy (1)


- Ideally programmers want memory that is
 - Large
 - Fast
 - Cheap
 - Nonvolatile
- Memory Hierarchy
 - Small amount of fast, expensive, volatile cache memory.
 - Some medium-speed, medium price, volatile main memory (RAM).
 - Gigabytes of slow, cheap, nonvolatile disk storage.
- Memory manager handles the memory hierarchy.

Storage Hierarchy (2)



Memory Manager

- Manage memory hierarchy
 - Monitor used and free memory.
 - Allocate memory to processes.
 - Reclaim (De-allocate) memory.
 - Swapping between main memory and disk.





Memory Allocation Schemes

- Contiguous Memory Allocation (3.1 & 3.2)
 - Each programs data and instructions are allocated a single contiguous space in memory.
 - Single Partition Allocation
 - Fixed-sized Partition Allocation
 - Variable-sized Partition Allocation
- Non-Contiguous Memory Allocation
 - Each programs data and instructions are allocated memory space that is not continuous.
 - E.g. Paged Memory Management, Segmented Memory Management

No Memory Abstraction

- Early mainframe, early minicomputers, early personal computers had no memory abstraction...
 - E.g. MOV REGISTER1, 1000
 - Here 1000 means move the content of physical memory address 1000 to register.
- No abstraction memory still used in embedded and smart card systems.
 - Radios
 - Washing machines
 - Microwave ovens
 - ...

Mono Program without Memory Abstraction (1)

- Three different ways to organize memory
 - (a) OS at the bottom of memory in RAM; was used in mainframes and minicomputers; not used anymore.
 - (b) OS is in ROM at the top of memory; is used in some palm-tops and embedded systems.
 - (c) Device drivers are in ROM at the top of memory and the rest is in RAM; was used by early PCs (e.g., running MS-DOS), where the portion in the ROM is called BIOS (Basic Input Output System).
 - Disadvantage of (a) and (c) : A bug in the program can wipe out the OS, possibly with disastrous results (such as garbling the disk)

Mono Program without Memory Abstraction (2)

- Three different ways to organize memory (ctd.)

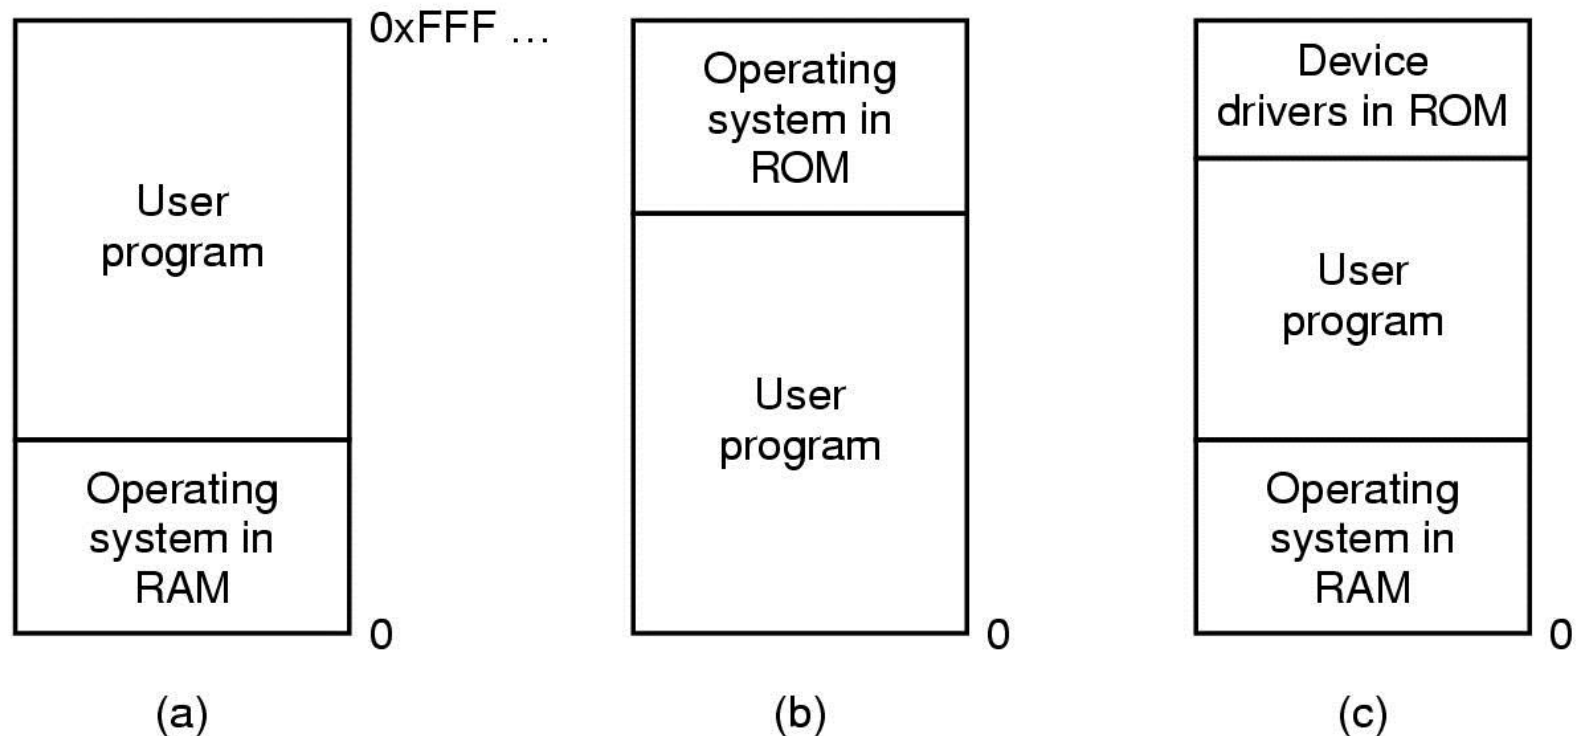


Figure 3-1. Three simple ways of organizing memory with an operating system and one user process.

Mono Program without Memory Abstraction (3)

- Single Partition Allocation.
- Note
 - A system using single contiguous allocation may still multitask by swapping the contents of memory to switch among users.
- Advantages
 - It is simple.
 - It is easy to understand and use.
- Disadvantages
 - It leads to poor utilization of processor and memory.
 - Users job is limited to the size of available memory.

Mono Program without Memory Abstraction (4)

- What to do when program size is larger than the amount of memory/partition (that exists or can be) allocated to it?
 - Overlays (past)
 - Dynamic Linking (Libraries – DLLs)(now)

Overlays (1)

- If a process is larger than the amount of memory, a technique called overlays can be used.
- Overlays is to keep in memory only those instructions and data that are needed at any given time.
- When other instructions are needed, they are loaded into space that was occupied previously by instructions that are no longer needed.
- Overlays are implemented by **programmer**, no special support needed from operating system, programming design of overlay structure is complex.

Overlays Example (1)

- Consider a two-pass assembler:
 - Pass1 constructs a symbol table.
 - Pass2 generates machine-language code.
 - Assume the following:

| | Size (k = 1024 bytes) |
|-------------------|-----------------------|
| Pass1 | 70k |
| Pass2 | 80k |
| Symbol table | 20k |
| Common routines | 30k |
| <i>Total size</i> | <i>200k</i> |

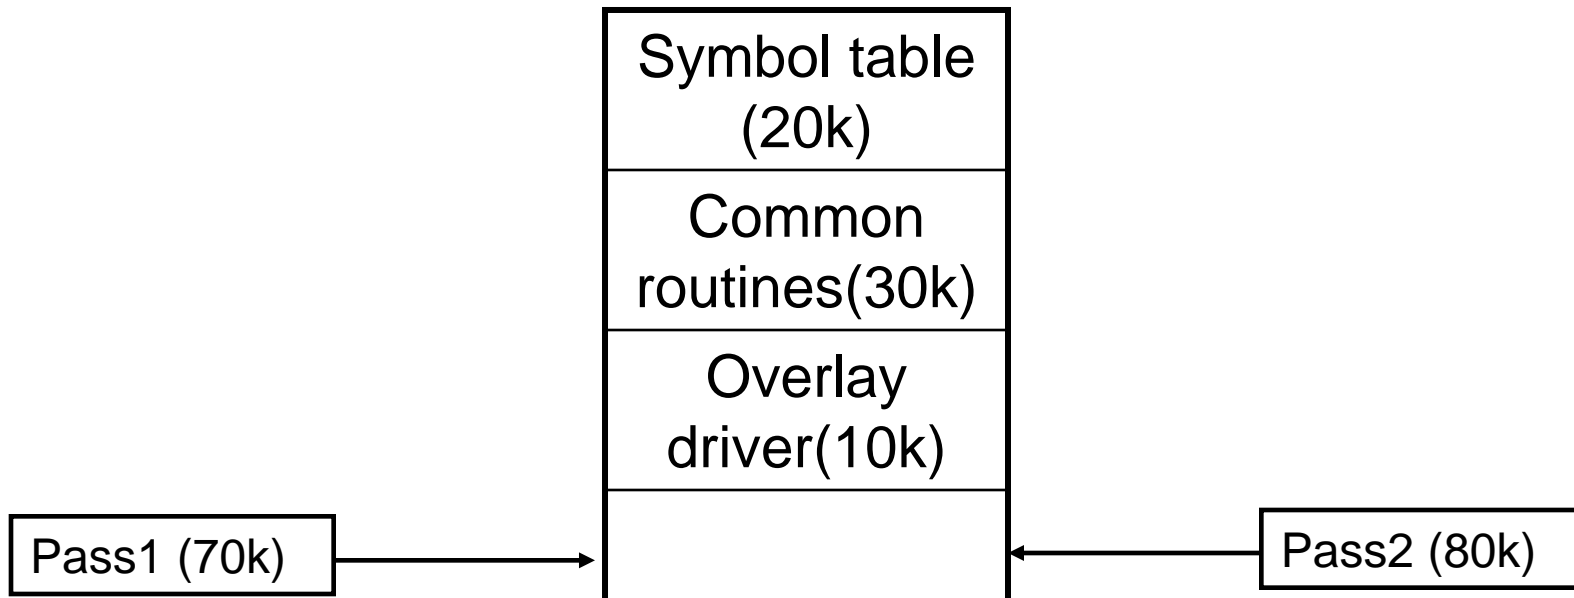
- To load everything at once, we need 200k of memory.

Overlays Example (2)

- Consider a two-pass assembler (ctd.)
 - If only 150K memory is available, we cannot run our process.
 - Note: **Pass1 and Pass2 do not need to be in memory at same time.**
 - Define two overlays:
 - Overlay A: symbol table, common routines, and Pass1.
 - Overlay B: symbol table, common routines, and Pass2.
 - Overlay Manager/Driver
 - Responsible for loading and unloading on overlay segment as per requirement.

Overlays Example (3)

- Consider a two-pass assembler (ctd.)
 - Add overlay driver 10k and start with overlay A in memory.
 - When finish Pass1, we jump to overlay driver, which reads overlay B into memory overwriting overlay A and transfer control to Pass2.
 - Overlay A needs 130k and Overlay B needs 140k.






Problems of Overlays

- Difficult for programmers to manage.
 - Requires programmers to organize overlay structure of program with the help of file structures etc.
 - Requires programmers to specify which overlay to load at different circumstances.
- Not possible to access overlays that are not loaded in the memory.


Fixed-Partition Multiprogramming (1)

■ Fixed-Partition Allocation

- Partition main memory into a set of non-overlapping regions called partitions.
- Each active process receives a fixed-size partition.
- Partitions fixed at boot time.
- Processor rapidly switches between each process.
- Partitions can be of equal or unequal sizes.
- Multiple boundary registers protect against damage.



Equal-size partitions



Unequal-size partitions

Fixed-Partition Multiprogramming (2)

■ Drawback of Fixed-Partition

- Internal fragmentation
- External fragmentation
- Limitation on the Size of the Process
- Degree of Multiprogramming is Less

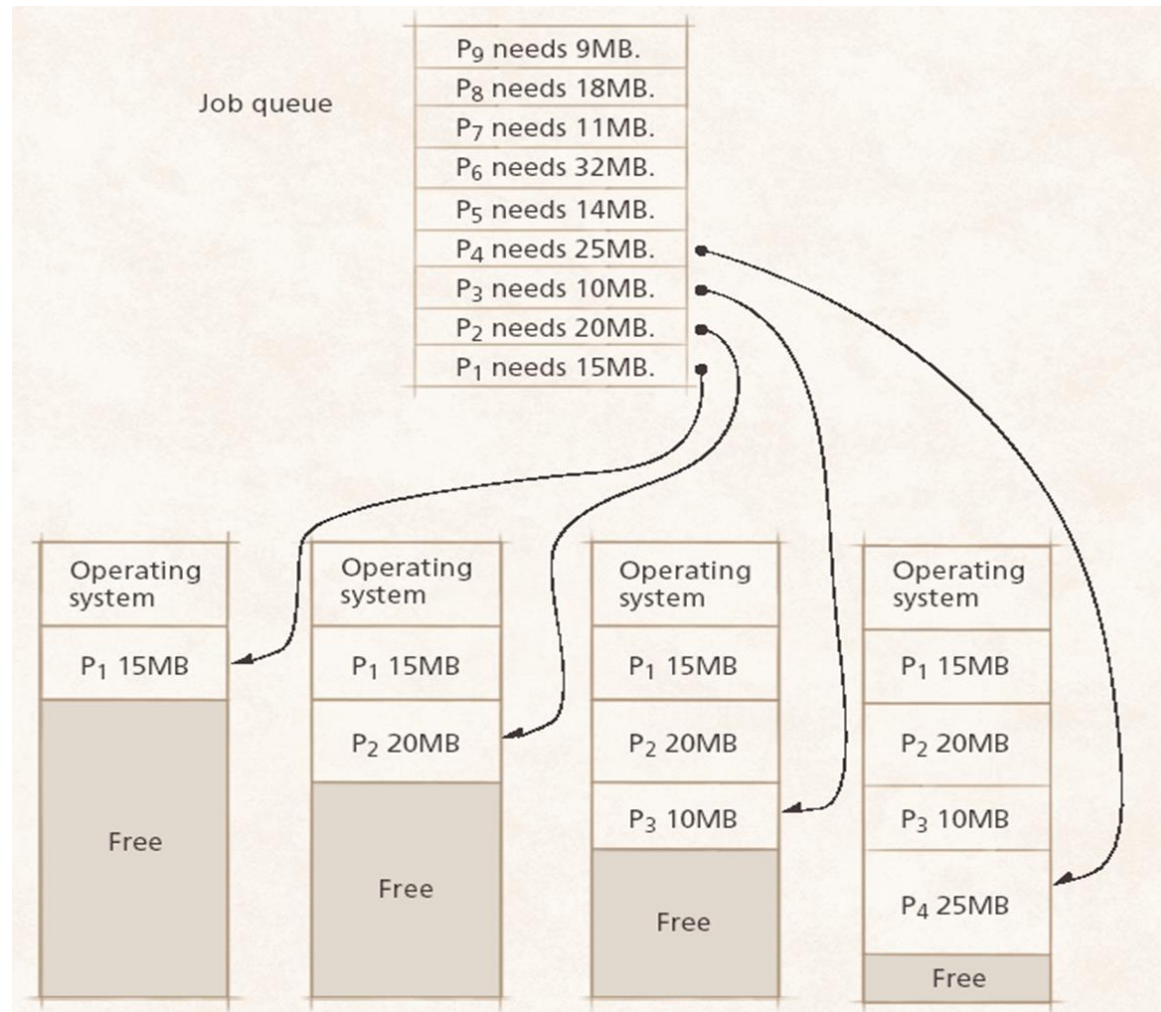
Variable-Partition Multiprogramming (1)

■ Variable -Partition Allocation

- Jobs placed where they fit.
 - No space wasted initially.
 - Internal fragmentation impossible
 - Partitions are exactly the size they need to be
- External fragmentation can occur when processes removed.
 - Leave holes too small for new processes.
 - Eventually no holes large enough for new processes.

Variable-Partition Multiprogramming (2)

■ Variable -Partition Allocation Example





Memory and Multiprogramming

- Memory needs two things for multiprogramming:
 - Relocation
 - Protection
- Relocation
 - Process may not be placed back in same main memory region!
 - How does a task or process run in different locations in main memory?
- Protection
 - How does the system prevent processes interfering with each other?
 - Must be ensured by processor (hardware) rather than OS.




Multiple Programs Without Memory Abstraction

- With Swapping Technique
- Without Swapping Technique

With Swapping

- Swapping is an approach for memory management by bringing each process in **entirely**, running it and then putting it back on the disk, so that another program may be loaded into that space.



Without Swapping (1)

- With the addition of some special hardware, it is possible to run multiple programs concurrently.
- Protection Solution of Early IBM 360
 - Divide main memory into 2KB blocks.
 - Assign each block a 4-bit protection key held in special registers inside the CPU.
 - E.g. A machine with a 1M memory, Need 512 4-bit registers → 256 bytes of key storage.
 - Each process also has a protection key value associated with it. (kept in PSW)
 - On a memory access the hardware checks that the current process's protection key matches the value associated with the memory block being accessed; if not, an exception occurs.

Without Swapping (2)

■ Relocation Problem of Early IBM 360

□ Example: two programs, each is 16KB.

- (a) 16 KB program runs alone. Finishes.
- (b) OS loads new 16KB program. Finishes
- (c) OS tries to run multiple programs. Instead of evicting old program, adds new one above it.

JMP 28 causes an error!

□ Problem

- Two programs both reference static, physical addresses.

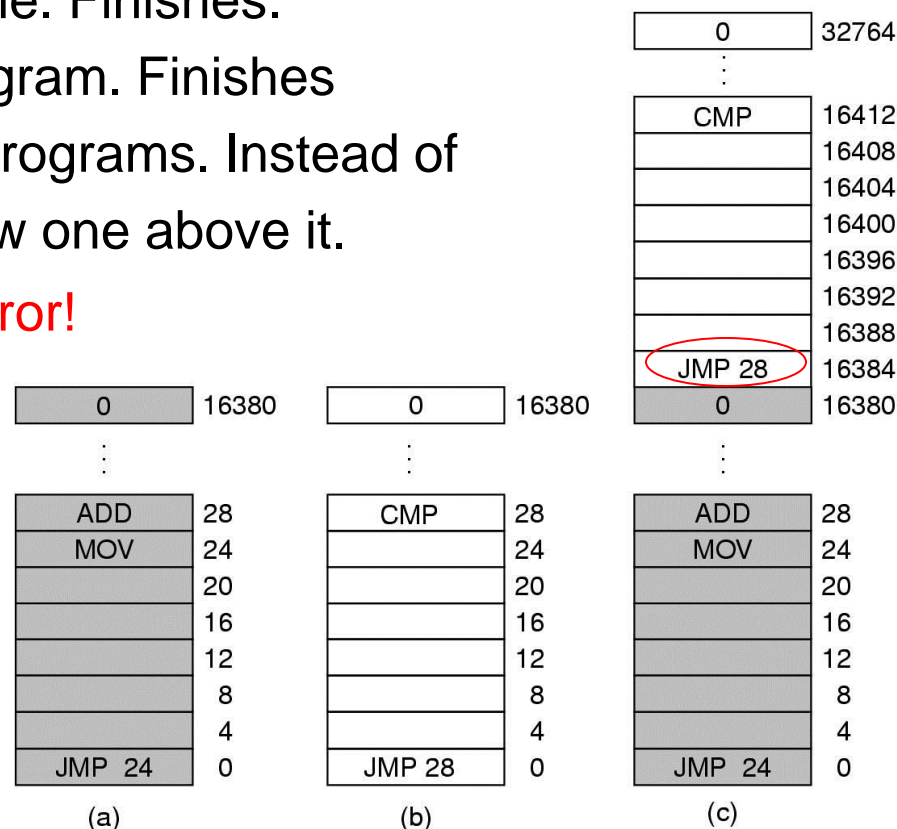
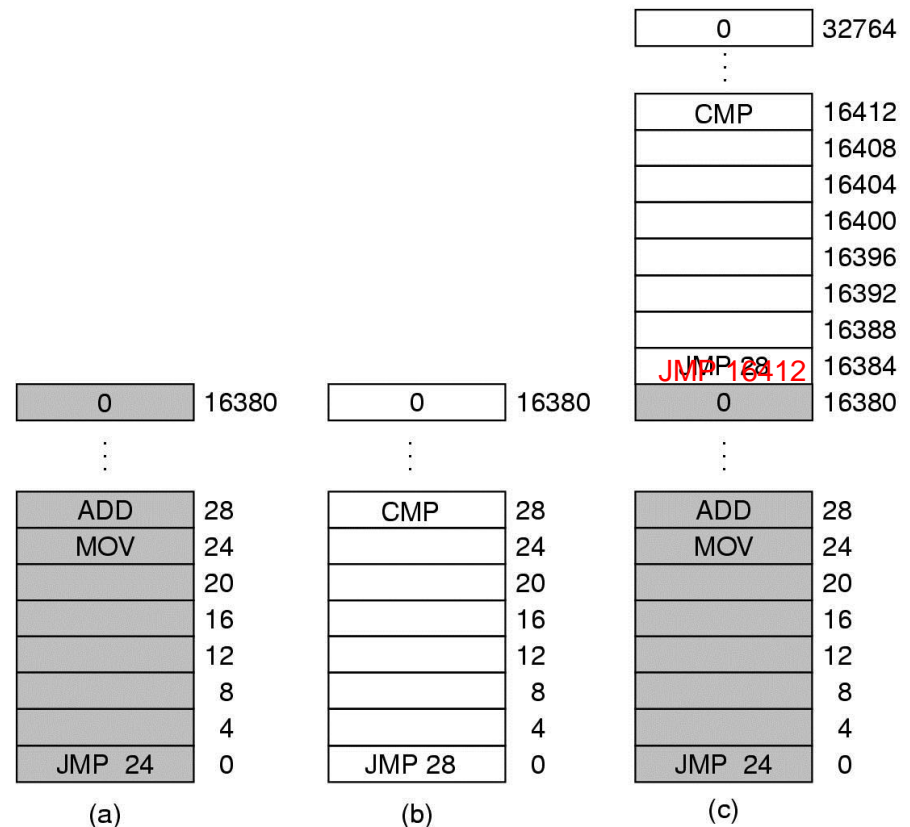


Figure 3-2. Illustration of the relocation problem.

Without Swapping (3)

■ Static Relocation

- Modify addresses statically (similar to linker) when load process.
- Example: two programs, each is 16KB.
 - $28 + 16384 = 16412$



Without Swapping (4)

■ Static Relocation (ctd.)

□ Advantage

- Requires no hardware support.

□ Disadvantages

- Slows down loading.
- Once loaded, the code or data of the program can't be moved into the memory without further relocation.
- The loader needs some way to tell what is an address and what is a constant.
 - E.g. `MOVE REGISTER1, 28`

data or address?





A Memory Abstraction: Address Space (1)

- Major drawbacks of exposing physical memory to process.
 - If user programs can address every byte of memory, they can easily trash the operating system.
 - It is difficult to have multiple programs running at once (take turns, if there is only one CPU).



A Memory Abstraction: Address Space (2)

- Recall: Memory needs two things for multiprogramming.
 - Relocation
 - How does a task or process run in different locations in main memory?
 - Protection
 - How does the system prevent processes interfering with each other?

A Memory Abstraction: Address Space (3)

- The Notion of an Address Space

- Address Space

- The set of addresses that a process can use to address memory (abstract memory for programs to live in).

- Each process has its own address space; independent of other processes.

- Address Binding

- The process of associating program instructions and data (addresses) to physical memory addresses is called address binding, or relocation.

A Memory Abstraction: Address Space (4)

■ Dynamic Relocation (Dynamic Binding)

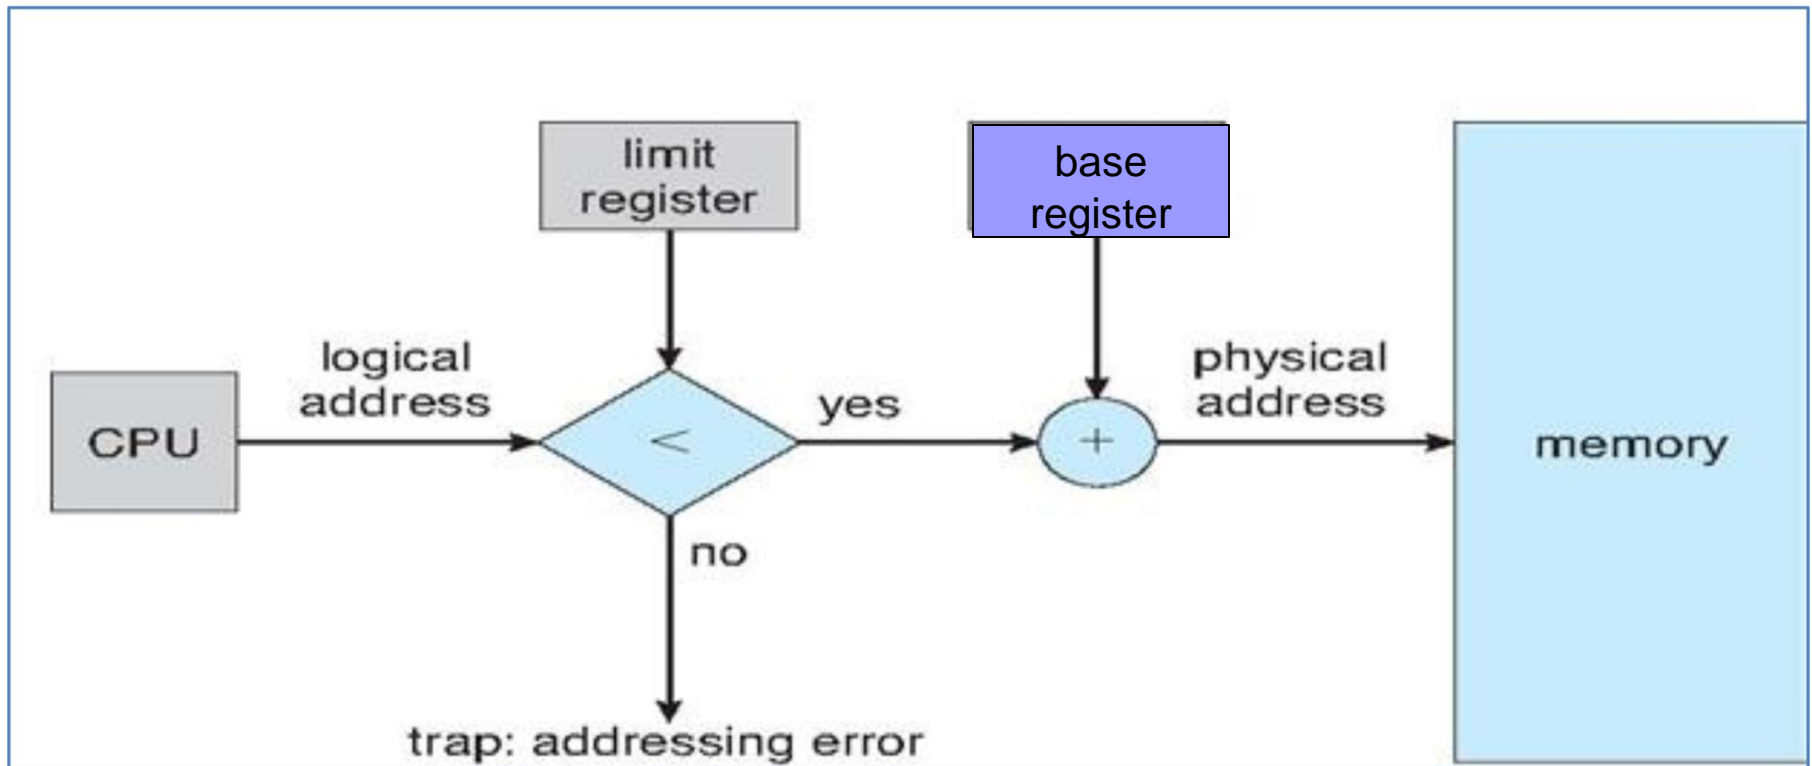
- Map each process's address space onto a different part of physical memory.

- Require Hardware Support

- Equip CPU with two special hardware registers: *base* and *limit*
- **Base register**: start location for address space
- **Limit register**: size limit of address space
- *These values are set when the process is loaded and when the process is swapped in.*

A Memory Abstraction: Address Space (5)

- Dynamic Relocation (ctd.)
 - Relocation Address Calculation



A Memory Abstraction: Address Space (6)

■ Dynamic Relocation (ctd.)

□ Example: two programs, each is 16KB.

■ First program

- Base register: 0
- Limit register: 16384

■ Second program

- Base register: 16384
- Limit register: 16384

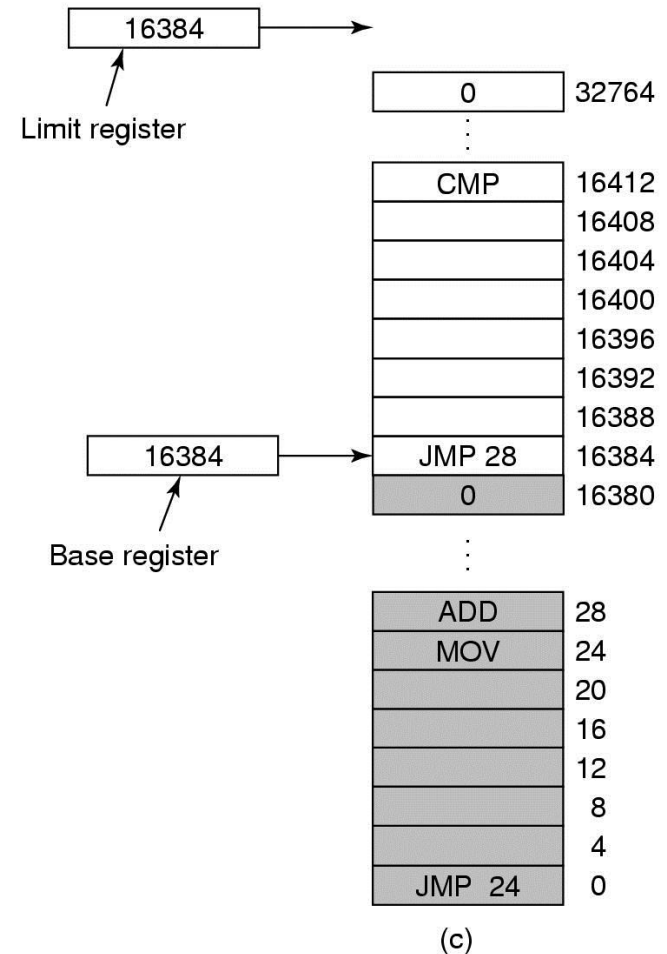


Figure 3-3. Base and limit registers can be used to give each process a separate address space.

A Memory Abstraction: Address Space (8)

■ Dynamic Relocation (ctd.)

□ Advantages

- OS can easily move process during execution.
- OS can allow process to grow over time.
 - OS just changes limit register or moves it.
- Simple, fast hardware
 - Two special registers, add, & compare

□ Disadvantages

- Slows everything (add on *every* reference).
- Can't share memory between processes.
- Process limited to physical memory size.
- Complicates memory management.

Address Binding (1)

■ Address binding of instructions and data to memory addresses can happen at three different stages.

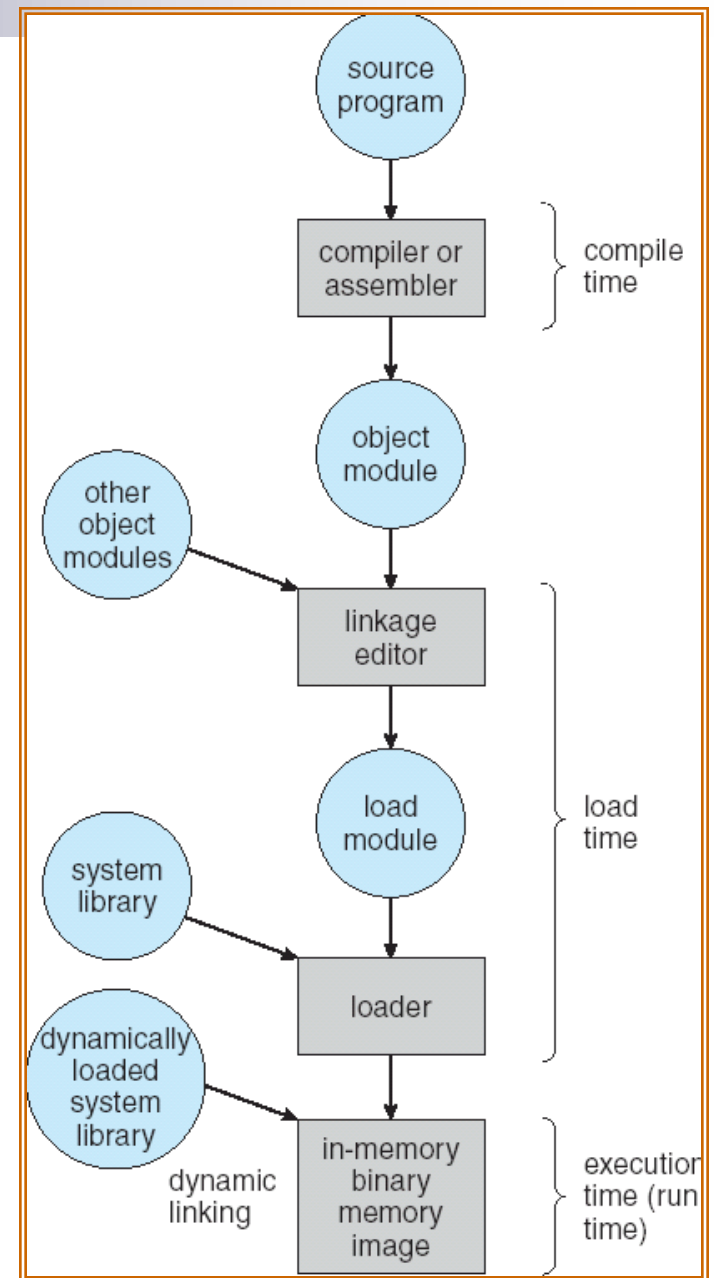
- Compile time
- Load time
- Execution time

Compiler → Generates Object Code

Linker → Combines the Object code into a single self sufficient *executable code*

Loading → Copies executable code into memory

Execution → dynamic memory allocation



Multi-step Processing of a User Program



Address Binding (2)

■ Compile-Time Binding

- Location of program in physical memory must be known at compile time.
- The compiler or assembler translates symbolic addresses (e.g., variables) to absolute addresses.
- Loading \equiv copying executable file to appropriate location in memory.
- If starting location changes, program will have to be recompiled.
- Example: .COM programs in MS-DOS

Address Binding (3)

■ Load-Time Binding

- Compiler generates relocatable code.
 - Compiler binds names to relative addresses (offsets from starting address).
 - Compiler also generates relocation table.
- Linker resolves external names and combines object files into one loadable module.
- (Linking) loader converts relative addresses to physical addresses.
- No relocation allowed during execution.



Address Binding (4)

■ Run-Time Binding

- Programs/compiled units may need to be relocated during execution.
- CPU generates relative addresses.
- Relative addresses bound to physical addresses at runtime based on location of translated units.
- Suitable hardware support required.



Two Relocation Policies (1)

Depending upon when and how the addresses translation from the logical address to physical address of main memory, takes place

■ Static Relocation

- Programs are loaded into consecutive memory locations wherever there is room and without relocation during loading.
- At load time, OS adjusts addresses in process to reflect position in memory.
- OS cannot move process after relocation.
- Example: OS/MFT



Two Relocation Policies (2)

■ Dynamic Relocation

- Run-time mapping of logical address into physical address with the support of some hardware mechanism.
- Hardware adds base register to logical address to get physical address.

Swapping (1)

- If there isn't room enough in memory for all processes, some processes can be swapped out to make room.
 - OS swaps a process out by storing its complete state to disk.
 - OS can reclaim space used by ready or blocked processes.
- When process becomes active again, OS must swap it back in (into memory)
 - With static relocation, the process must be replaced in the same location.
 - With dynamic relocation, OS can place the process in any free partition (must update the base and limit registers).
- Swapping and dynamic relocation make it easy to increase the size of a process and to compact memory (although slow!)

Swapping (2)

■ Example




Figure 3-4. Memory allocation changes as processes come into memory and leave it. The shaded regions are unused memory.

Swapping (3)

■ Context Switch

- Swapped in & out cost too much.
- Assume: process size 1MB, disk transfer rate 5MB/sec, average latency 8ms
 - Transfer time = $1\text{MB} / (5\text{MB/sec}) = 1/5 \text{ sec} = 200 \text{ ms}$
 - Swap time = 208 ms
 - Swap out & in = 416 ms
- Major part of swap time is transfer time.
- For RR scheduling, time quantum should $\gg 416\text{ms}$
- Modified versions of swapping are found on many systems (i.e., UNIX, Linux, and Windows).

Swapping (4)

- How to allocate when a process is created or swapped in?
 - If process are created with a fixed size, OS allocates exactly what is needed.
 - What if most processes will grow as they run?
 - E.g. one growing segment: data segment
 - E.g. two growing segments: stack segment and data segment

Swapping (5)

■ Solution for Growing Process

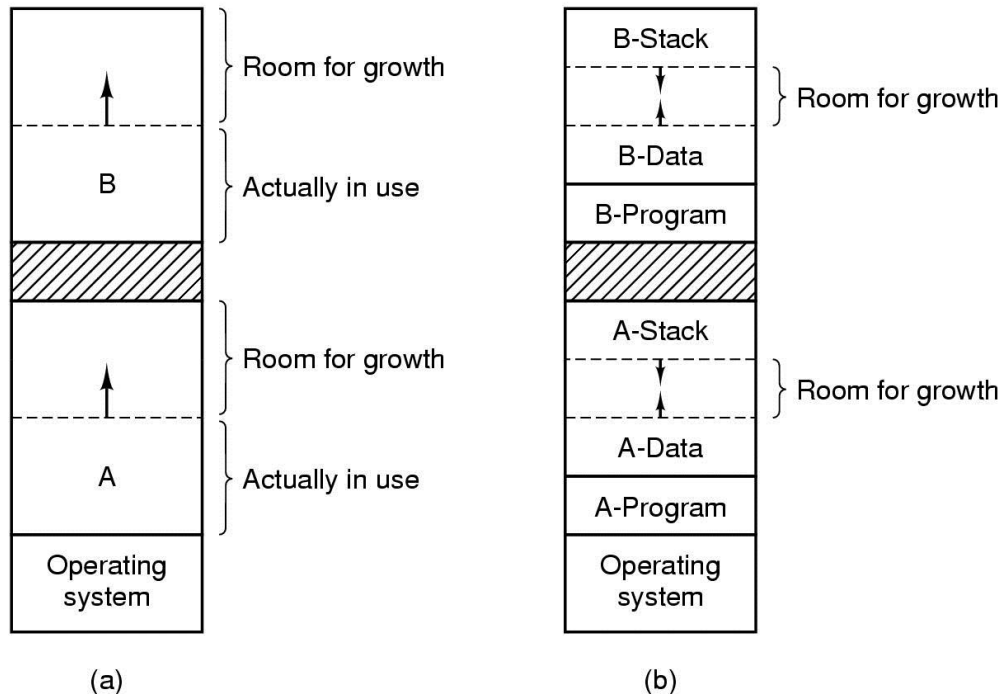


Figure 3-5. (a) Allocating space for growing data segment.
(b) Allocating space for growing stack, growing data segment.



Problems of Swapping (1)

■ Problem 1

- Swapping creates holes in memory (external fragmentation).

- Can be solved using compaction.

■ Compaction

- Memory contents shuffled to place all free memory together in one large block.

- Dynamic relocation (run-time binding) needed.

Problems of Swapping (2)

■ Compaction Example

- Assumes programs are all relocatable
- Processes must be suspended during compaction
- Need be done only when fragmentation gets very bad





Problems of Swapping (3)

■ Problem 2

- Process must fit into physical memory (impossible to run larger processes).
 - Can be solved using overlays.




Managing Free Memory

- How to keep track of memory usage?
 - Bitmaps and free lists.
- Bitmap
 - Dividing memory into small units; corresponding to each unit is a bit in the bitmap; 0 if the unit is free, 1 if units is occupied.
 - When allocate, memory manager has to find consecutive 0 bits.
- Linked Lists
 - Each entry in the list specifies a hole or a process, the address it starts, the length, and a pointer to the next entry.

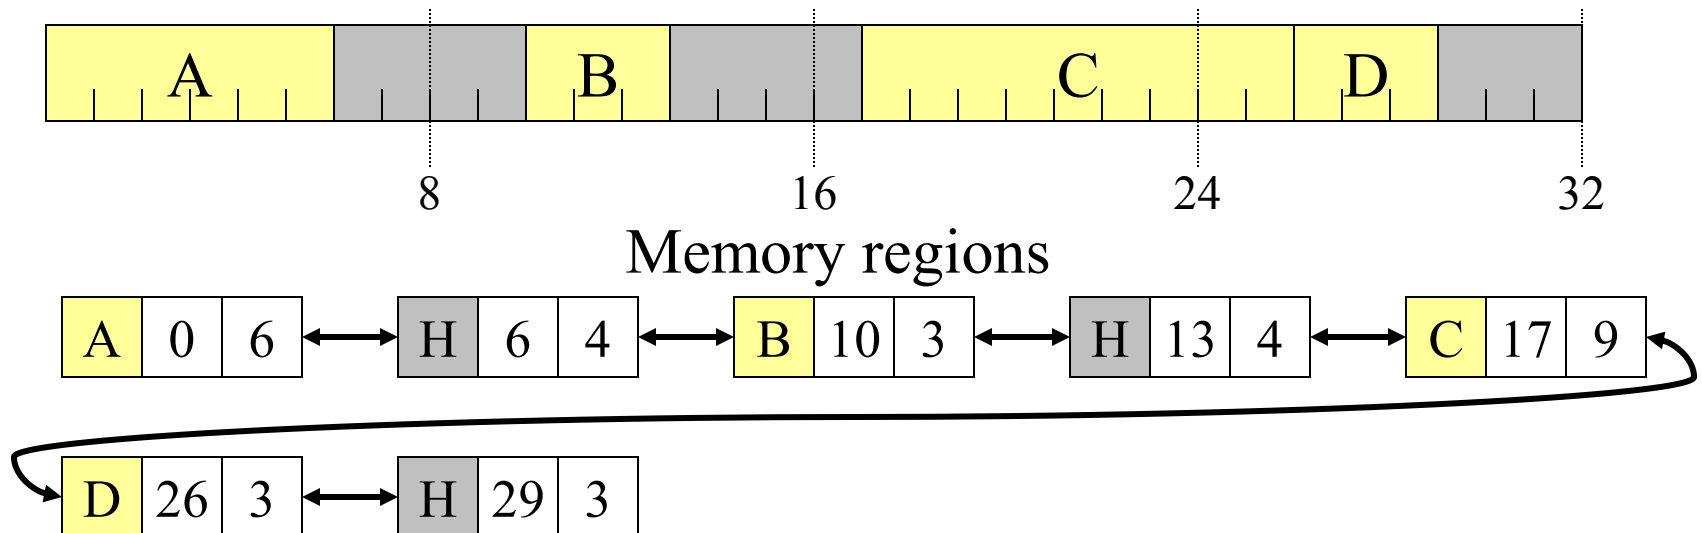
Memory Management with Bitmaps

- Memory is divided into allocation units.
 - One allocation unit corresponds to 1bit in the bitmap
 - 0: free, 1: allocated
- Allocation unit size determines efficiency.
 - At 1 bit per 4KB chunk, we need just 256 bits (32 bytes) per MB of memory.
 - For smaller chunks, we need more memory for the bitmap.
 - Can be difficult to find large contiguous free areas in bitmap.




Memory Management with Linked Lists

- (1)
- Keep track of free/allocated memory regions with a linked list.
 - Each entry in the list corresponds to a contiguous region of memory.
 - Entry can indicate either allocated or free (and, optionally, owning process).
 - May have separate lists for free and allocated areas.
 - Efficient if allocation unit size is large.
 - Fixed-size representation for each region.
 - More regions => more space needed for free lists.



Memory Management with Linked Lists (2)

- Four neighbor combinations for the terminating process X



Storage Placement Strategies (1)

How to satisfy a request of size n from a list of free holes?

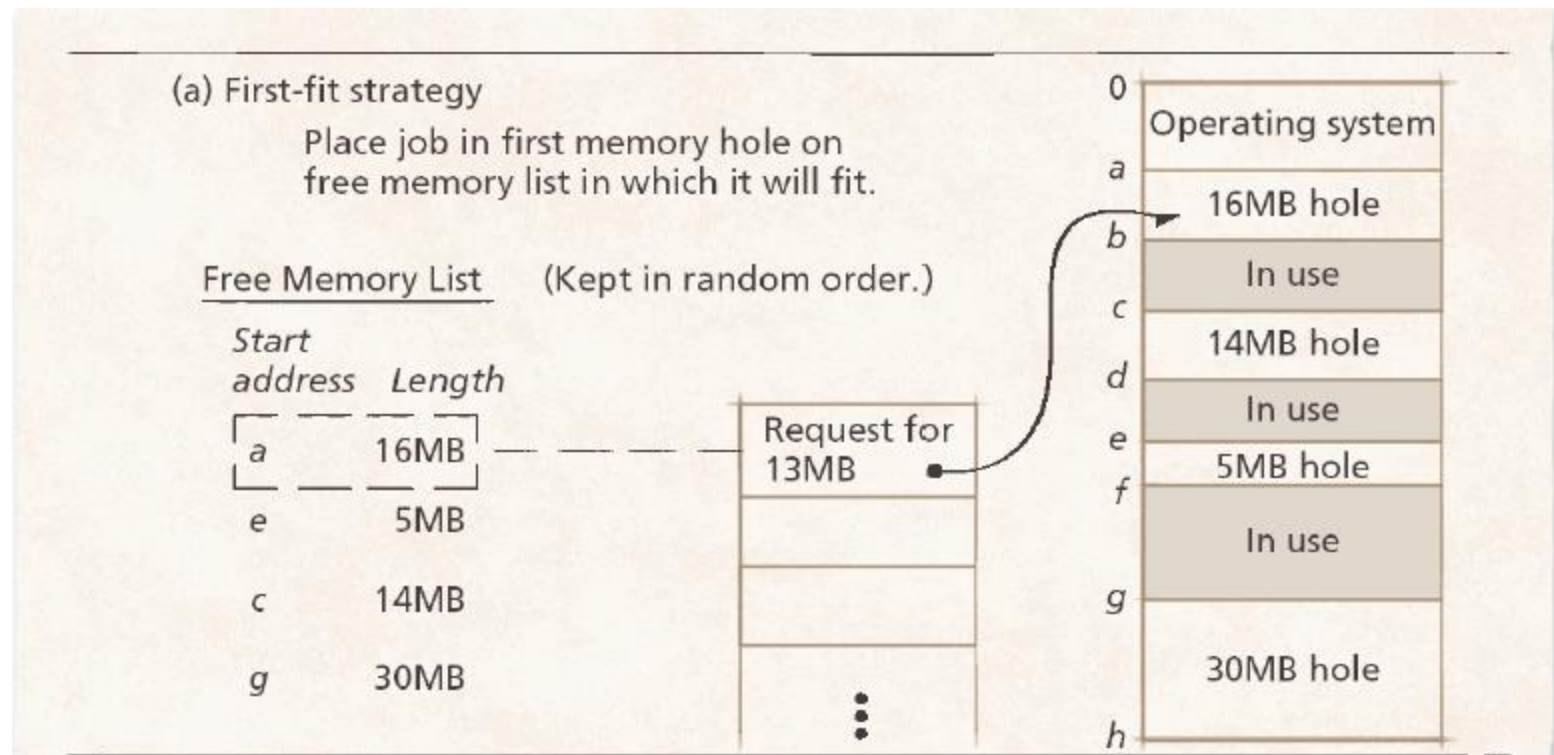
■ First Fit

- Scan the list, use the first available hole whose size is sufficient to meet the need.
 - If larger in size, break it into an allocated and a free part.
- May have many processes loaded in the front end of memory that must be searched over when trying to find a free block (a bit inefficient).
- May have lots of unusable holes at the beginning.
 - External fragmentation

Storage Placement Strategies (2)

■ First Fit (ctd.)

□ Example



Storage Placement Strategies (3)

■ Next Fit

- Minor variation of first fit.
- It begins its search from that point in the list where the last request succeeded.
- Problem
 - Slightly worse performance than first fit.

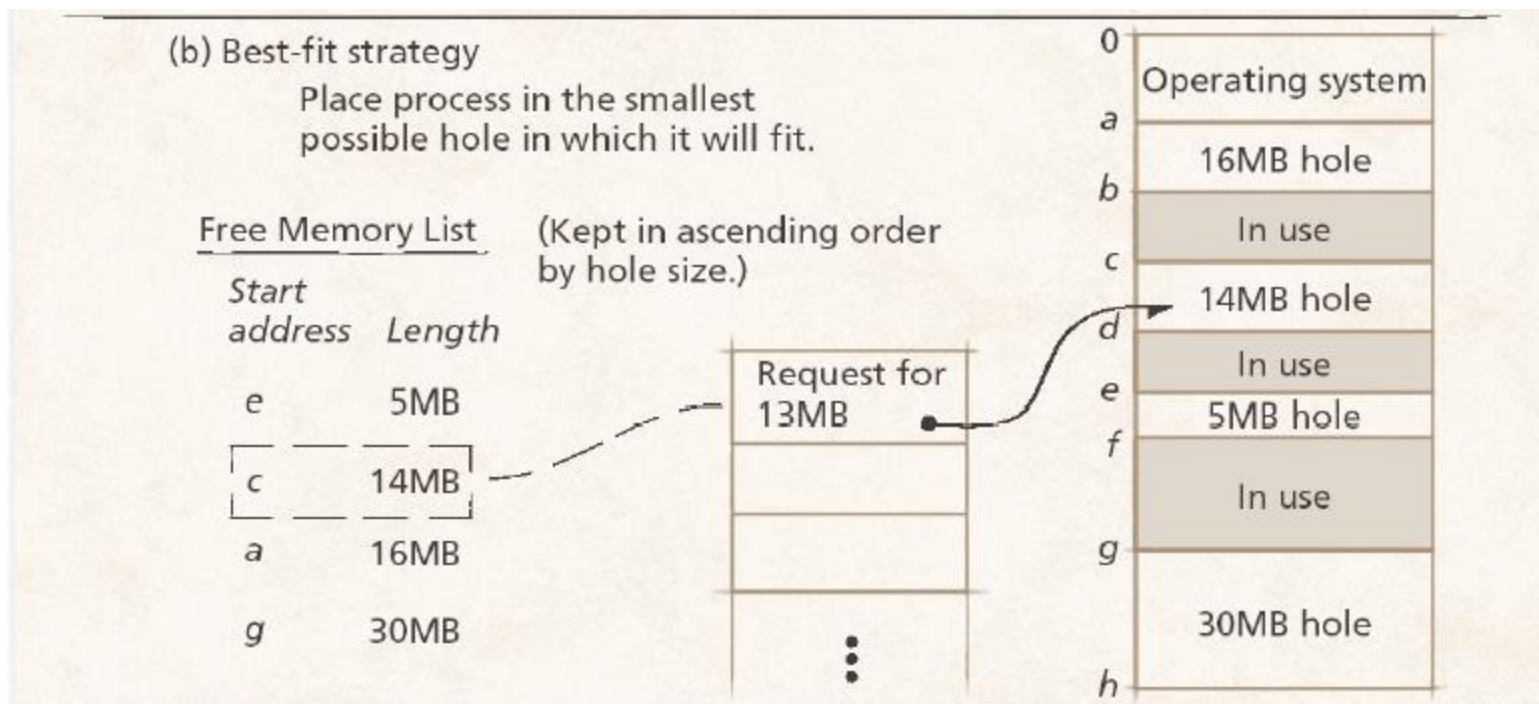
■ Best Fit

- Use the hole whose size is equal to the need, or if none is equal, the hole that is larger but closest in size.
- Problem
 - Often have to search the complete list.
 - Creates small holes that can't be used.

Storage Placement Strategies (4)

■ Best Fit (ctd.)

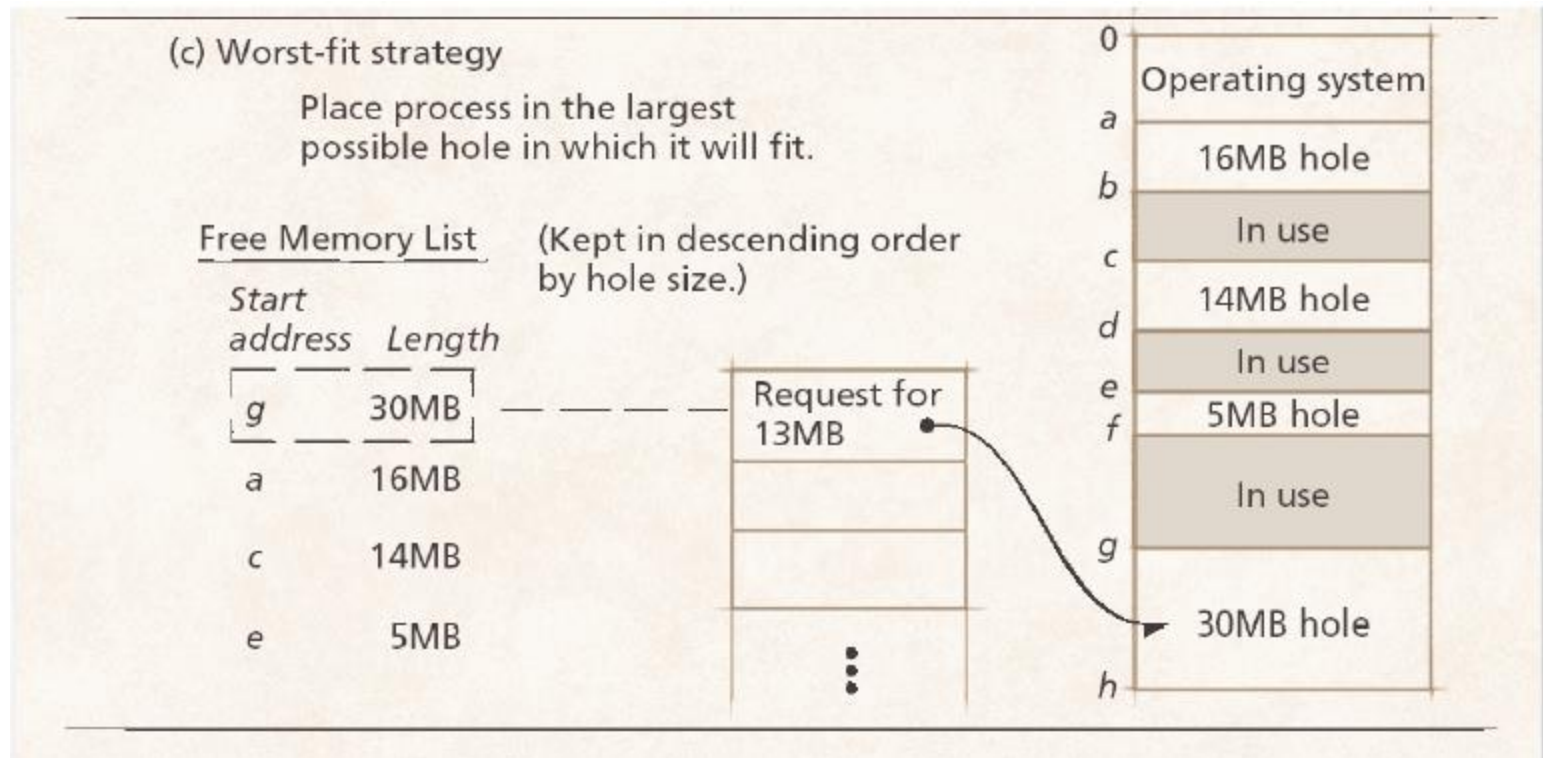
□ Example



Storage Placement Strategies (5)

■ Worst Fit

- Use the largest available hole.
- Problem
 - Often has to search complete list.
 - Gets rid of large holes making it difficult to run large programs.



Storage Placement Strategies (6)

■ Quick Fit

- Maintains separate lists for some of the more common sizes requested.
- When a request comes for placement it finds the closest fit.
- This is a very fast scheme, but a merge is expensive. If merge is not done, memory will quickly fragment in a large number of holes into which no processes fit.



Summary

- Memory Allocation
 - Contiguous Allocation
 - Non-contiguous Allocation
- Contiguous Allocation
 - Relocation
 - Static Relocation
 - Dynamic Relocation
 - Overlays
 - Swapping
- Managing Free Memory
 - Bitmap
 - Link list
- Storage Placement Strategies
 - First fit
 - Best fit
 - Worst fit



Homework

■ P254 3, 4